

Applied Machine Learning

Text Retrieval, Association Rules and Decision Tree Learning using R

Adrian Cuyugan / Information Analytics, GAMBI

Contents

1	Introduction	2
2	Scope	2
2.1	Techniques	2
2.2	Before starting	2
3	Data	3
3.1	Source and Goal	3
3.2	Bag of Words	3
3.3	Data fields	3
4	Text Retrieval	3
4.1	Loading the dataset	3
4.2	Building the corpus	4
4.3	Sparsity	4
4.4	Combining the term matrix with the response variable	5
5	Association Rules	5
5.1	Negative reviews	5
5.2	Random sampling	5
5.3	A Priori algorithm	6
5.4	Pruning	7
5.5	Custom stopwords	8
5.6	Visualization	8
6	Decision Tree Learning	11
6.1	Most frequent words	11
6.2	Bit-vector weighting	11
6.3	Divide train and testing set	11
6.4	Build tree	11
6.5	Visualize tree	12
6.6	Complexity parameter	12
6.7	Model evaluation	14

1 Introduction

Usually, Machine Learning stops after you understand how the algorithm runs. Data Mining, on the other, is the real-life application of how these complex algorithms is put into practice with very noisy data - recognizing any patterns, discovering outliers and estimating coefficients which and can later be developed to produce accurate predictions.

2 Scope

2.1 Techniques

This document assumes that you have basic knowledge on the following subjects:

1. **Information Retrieval** - the lesson stopped at text retrieval since the data used is always offline and does not necessarily require a constant query function, $f(q, d)$. For this activity, we will be using [tm](#) package to create a corpus and build the sparse matrix. We will be using bag of words to construct our sparse matrix using the raw term frequency count of the vector space model.

$$VSM(q, d) = \sum_{i=1}^n x_1 y_1 \dots + \dots x_n y_n$$

2. **Association Rules** - finding frequent patterns based on the text corpus we have built from [tm](#). [arules](#) and [arulesViz](#) will be used for this. Also, [Gephi](#) was used to visualize the network of rules created from the Apriori algorithm. *Corpus* is a latin word which is a collection of plain text that can be read by the computer when doing Natural Language Processing - linguistics, building sparse matrix, frequency count, etc.
3. **Decision Tree** - the simplest Classification and Regression Tree (CART) algorithm computes for entropy based the response variable that we set, *sentiment*. Remember that entropy is:

$$E = - \sum_{i=1}^k \frac{C_i}{n} \log_2 \frac{C_i}{n}$$

We will be using [rpart](#) to estimate the model and will be using [rattle](#) to visualize the tree.

2.2 Before starting

You can use `?help` to search what are the description and available parameters that you can use in this example. For example, if you want to know what does `head()` do, type `?head` in R console.

I made sure that scripts were coded as simple as possible and have attached comments so that you can understand what happens within each line of code.

The scripts do not include `install.packages()` as it assumes that you have already installed the packages. If you have not, please refer to the above techniques as the names of packages are highlighted with their corresponding documentation from CRAN directory.

3 Data

3.1 Source and Goal

A current knowledge-based competition (bragging rights award) is being held at [Kaggle](#). One of the available data set we could is [Google's Word2Vec's](#). For this exercise, we will not be performing analysis using Natural Language Processing (parsing of words) but only be doing bag of words.

Since prediction was not discussed (though I highly suggest that you make further reading on this), we will just be describing the estimates of the model. There is a big difference between building the model and interpreting the results than doing a prediction that leads to prescribing recommendation as these models are tested against another source of data set which is usually the testing set and/or cross-validation set.

3.2 Bag of Words

As discussed, *bag of words* will be used to create the vector of words based on all 25,000 movie reviews. These are also labelled so that we can calculate the estimates of probability of a word producing a positive or negative review. Sampling was also done to produce the association rules as it requires more memory to run the algorithm.

3.3 Data fields

1. **id** - Unique ID of each review. This will be excluded from the dataset but is very helpful in identifying the unique records of review. `rowname()` function in R accomplishes this task.
2. **sentiment** - The response variable. This is the sentiment of the review; 1 for positive reviews and 0 for negative reviews. We will be dropping this when we're doing unsupervised learning.
3. **review** - This is a free-form text field where the users wrote their reviews. This will be converted in to a *Document Term Matrix* as a result of the text retrieval process.

4 Text Retrieval

4.1 Loading the dataset

The CSV file is downloaded for offline use so that this document can be generated faster. The data set can be saved into the working directory. You can check your current working using `getwd()` or manually set it using `setwd()`. I used my C drive to save my data for easier viewing of the codes.

```
train <- read.csv("C://popcorn-train.csv", head = T, sep = ",") # Load data set

rownames(train) <- train$id # Assigns the ID as row identifier
train <- train[,2:3] # Subsets only Sentiment and Comments
train$sentiment <- factor(train$sentiment,
                          levels = c(1,0),
                          labels = c("Pos", "Neg")) # Converts 0 and 1 into categories
table(train$sentiment) # Count of records per sentiment
```

```
##
##   Pos   Neg
## 12500 12500
```

After loading the data set into R environment, I assigned the **id** as the row identifier then, I dropped the ID variable since it has already been attached. Since R automatically recognizes **sentiment** as a numerical variable, I changed it to categorical by using the **factor()** function.

Note that the proportion of Positive and Negative reviews are equal.

4.2 Building the corpus

The corpus can be built from a data frame or a vector format. This example uses data frame since the CSV have been loaded as a data frame.

Each line of code converts the corpus into its respective transformation function. For more advanced control of parameter, please refer to `?tm_map`.

```
# Prepare corpus
require(tm)
keyword <- as.data.frame(train$review)
keyword <- Corpus(DataframeSource(keyword)) # Creates corpus
keyword <- tm_map(keyword, removeWords, stopwords(kind = "english")) # Removes stop words
keyword <- tm_map(keyword, stemDocument, language = "porter") # Applies porter stemming
keyword <- tm_map(keyword, stripWhitespace) # Removes whitespace
keyword <- tm_map(keyword, content_transformer(tolower)) # Converts to lowercase
keyword <- tm_map(keyword, removePunctuation) # Removes punctuation marks
keyword <- tm_map(keyword, removeNumbers) # Removes numbers
```

```
# Raw Term Frequency
keyword.term <- DocumentTermMatrix(keyword) # Creates a sparse matrix
keyword.term # Checks for how many terms were produced
```

```
## <<DocumentTermMatrix (documents: 25000, terms: 102984)>>
## Non-/sparse entries: 2459554/2572140446
## Sparsity           : 100%
## Maximal term length: 60
## Weighting           : term frequency (tf)
```

There are more than 100,000 terms produced by the algorithm. If we are doing analysis on Information Retrieval, we would want all of these terms included but since we just need those that have effect on the model, we can drop the sparsed terms that barely have frequency count in the matrix.

4.3 Sparsity

Use `removeSparseTerms` to remove low frequency count terms.

```
keyword.term <- removeSparseTerms(keyword.term, sparse = 0.9)
keyword.term # Checks for sparsity
```

```
## <<DocumentTermMatrix (documents: 25000, terms: 126)>>
## Non-/sparse entries: 606899/2543101
## Sparsity           : 81%
## Maximal term length: 8
## Weighting           : term frequency (tf)
```

We can minimize the non-essential terms to almost 130. Although this is still a lot of variables that CART would handle, based on a data mining analyst's judgement, this can still be adjusted to < 0.90 which would lessen the terms included in the final model.

4.4 Combining the term matrix with the response variable

As CART requires a response variable, we will be combining the document term matrix with the sentiment vector.

```
keyword.matrix <- as.data.frame.matrix(keyword.term) # Converts to data frame matrix
keyword.matrix <- cbind(keyword.matrix, sentiment = train$sentiment) # Combines response
rownames(keyword.matrix) <- row.names(train) # Applies row identifier
head(keyword.matrix[150:170,118:127]) # Previews the combined data frame
```

```
##          what whole will without wonder work world year young sentiment
## 11973_4      0      0      0          0          0      0      0      0      0      Neg
## 4541_10      0      0      0          0          0      0      0      0      0      Pos
## 5169_7       0      1      0          0          0      0      0      0      0      Pos
## 6809_1       0      0      0          0          0      0      0      1      0      Neg
## 8125_7       0      0      0          0          0      3      0      1      0      Pos
## 4802_8       0      0      0          0          0      1      0      0      1      Pos
```

5 Association Rules

5.1 Negative reviews

For this activity, I am only interested in looking at the bag of words from the negative reviews. Note that when we built the matrix, we used the raw term frequency. If we used the TF-IDF weighting, the values of the matrix would be totally different.

5.2 Random sampling

Building models take a lot of memory resources and when you hit the ceiling, your computer will definitely crash. To minimize the computational load, I sampled the negative reviews down to 5,000 records.

```
popcorn.neg <- subset(keyword.matrix, sentiment == "Neg") # Subsets negative reviews
popcorn.neg <- popcorn.neg[, !(colnames(popcorn.neg) %in% "sentiment")] # Drops response
# Random sampling
set.seed(14344) # For reproducibility
popcorn.neg <- popcorn.neg[sample(1:nrow(popcorn.neg), 5000, replace=FALSE),] # Samples
popcorn.neg <- as.matrix(popcorn.neg) # Converts data frame to matrix
```

For the results to be fully reproducible, `set.seed()` was used so that when you generate your own resampling method, we would have the same records.

5.3 A Priori algorithm

One of the most common algorithms in doing association rules is a priori.

Before running the function, the data set needs to be converted from matrix/data frame into a transactions class. For example:

ID	Transactions/Bag of Words
1	example, film, interested, just, like, look, movie, scene, will
2	because, beyond, director, movie, know, like, movie, want
3	because, better, forget, get, know, look, movie, not, will, work
4	better, interested, movie, same, something, work

```
require(arules)
popcorn.neg <- as(object = popcorn.neg, Class = "transactions") # Converts to transaction
popcorn.neg
```

```
## transactions in sparse format with
## 5000 transactions (rows) and
## 126 items (columns)
```

After sampling, we would have a reproducible random sample of 5,000 records with 127 bag of words.

```
neg.arules <- apriori(popcorn.neg, parameter=list(support=0.1, conf=0.1)) # Runs algorithm
```

```
##
## Parameter specification:
## confidence minval smax arem aval originalSupport support minlen maxlen
##      0.1      0.1      1 none FALSE          TRUE      0.1      1      10
## target  ext
## rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## apriori - find association rules with the apriori algorithm
## version 4.21 (2004.05.09)      (c) 1996-2004  Christian Borgelt
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[126 item(s), 5000 transaction(s)] done [0.01s].
## sorting and recoding items ... [113 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 done [0.01s].
## writing ... [3099 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

```
neg.arules <- sort(neg.arules, decreasing=T, by="lift") # Sorts based on lift values
```

The *minsup* (relative support) and *minconf* (conditional probability) values totally depend on the analyst's judgement. Also, expertise in the domain helps. Confidence is computed as follows:

$$conf(X \Rightarrow Y) = \frac{sup(X \cup Y)}{sup(x)} = P(Y|X) = \frac{P(Y \cap X)}{P(X)}$$

Find time to look at the output of the model. Based on the $minsup = 10\%$ and $minconf = 10\%$, there were 3,099 rules that were discovered. The $minlen$ can also be tweaked.

Then, the rules are sorted based on the magnitude of lift values. Lift is computed:

$$lift(X, Y) = \frac{conf(X \Rightarrow Y)}{sup(Y)} = \frac{sup(X \cup Y)}{sup(X) * sup(Y)}$$

5.4 Pruning

One property of frequent patterns is the downward closure where an itemset subset is also frequent if the parent itemset is also frequent. Pruning resolves this to remove ambiguity in the rules.

```
neg.prune <- is.subset(x = neg.arules, y = neg.arules) # Finds subsets
neg.prune[lower.tri(x = neg.prune, diag = TRUE)] <- NA # Discards lower triangle of matrix
neg.redundant <- colSums(neg.prune, na.rm = TRUE) >= 1 # Computes for sums of terms
neg.prune <- neg.arules[!neg.redundant] # Removes redundancies based on subsets

# Converts rules into data frame
# elements is a sub-class of the rule, just like subsetting
neg.arules.df = data.frame(
  lhs = labels(lhs(neg.prune))$elements,
  rhs = labels(rhs(neg.prune))$elements,
  neg.prune@quality)

head(neg.arules.df, 20) # Views top 20 rules based on lift
```

##	lhs	rhs	support	confidence	lift
## 2758	{just,like,movi}	{movie}	0.1068	0.5297619	1.576672
## 2726	{just,like,the}	{look}	0.1038	0.4556629	1.546717
## 1834	{movi,watch}	{movie}	0.1272	0.5170732	1.538908
## 2754	{even,movi,the}	{movie}	0.1044	0.5158103	1.535150
## 2766	{just,movi,the}	{movie}	0.1182	0.5025510	1.495688
## 1726	{like,make}	{look}	0.1084	0.4395783	1.492119
## 2770	{like,movi,the}	{movie}	0.1320	0.4988662	1.484721
## 1846	{good,movi}	{movie}	0.1244	0.4987971	1.484515
## 1873	{even,movi}	{movie}	0.1276	0.4984375	1.483445
## 1858	{get,movi}	{movie}	0.1214	0.4979491	1.481991
## 1825	{can,movi}	{movie}	0.1014	0.4970588	1.479342
## 2738	{just,like,the}	{realli}	0.1016	0.4460053	1.478797
## 1840	{movi,see}	{movie}	0.1218	0.4963325	1.477180
## 1828	{bad,movi}	{movie}	0.1206	0.4930499	1.467410
## 1852	{movi,time}	{movie}	0.1208	0.4918567	1.463859
## 1738	{just,like}	{look}	0.1246	0.4281787	1.453424
## 1900	{like,movi}	{movie}	0.1696	0.4870764	1.449632
## 2994	{just,make,the}	{even}	0.1020	0.5723906	1.447624
## 1891	{just,movi}	{movie}	0.1512	0.4861736	1.446945
## 482	{want}	{see}	0.1184	0.5262222	1.435412

```
neg.prune # Total rules after pruning
```

```
## set of 1249 rules
```

You can use `inspect()` to view the rules from the `arules` package. Unfortunately, since we loaded `tm` package in the environment and `tm` uses this function for a different purpose, we cannot use `inspect()`. That's why I had to convert the rules into a data frame so that it can be easily exported or viewed.

There were a total of 1,249 rules after pruning.

Going back to the top 20 rules, assume interpretations of patterns found by the a priori algorithm.

5.5 Custom stopwords

Now, the goal is to construct association based on the rules found. As the rules that merit high lift values do not make sense, *just*, *movie*, *really*, these can be added to a custom stop words. Then, another run of pattern mining could be done to find more meaningful words. For this exercise, let's assume that these words make sense to construct meaningful assumptions or **if** \Rightarrow **then** statements.

Or another approach is to use TF-IDF because common words such as *movi* and *movie* score low using this weighting. Remember that a normalized TF-IDF is computed as:

$$IDF_{TF} = \log\left[\frac{(M + 1)}{k}\right]$$

5.6 Visualization

The first 40 rules ranked by the magnitude of the lift value are visualized to show their association. The size of the bubbles is the *support* and the color density of it is the *lift*.

Precision is set to show tenth decimal and `cex` is a numerical value giving the amount by which plotting text and symbols should be magnified relative to the default.

```
require(arulesViz)
par(mfrow=c(1,2)) # Displays two plots in one graphics
plot(head(neg.prune,40), method="graph",
      control=list(type="items", precision=1, cex=.6,
                  main="Negative Reviews Items")) # Plots items rules
plot(head(neg.prune,40), method="graph",
      control=list(type="itemsets", precision=1, cex=.6,
                  main="Negative Reviews Itemsets")) # Plots itemsets
```

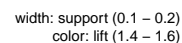
```
par(mfrow=c(1,1)) # Resets graphic device
```

Using some basic understanding in graph theory (which is another topic - centrality, degrees, clustering, etc), we can cluster these words based on their association with one another. I would let you interpret the graphs for deeper understanding.

Another approach is to perform community modularity clustering based on all rules found in the model. [Gephi](#) is particularly effective in this and it is open-source. Figure 2 is an example.

We can also perform basic NLP by just looking for patterns using *syntactic parsing* (nouns vs verbs) compared with positive reviews. By parsing the bag of words, more advanced techniques such as *Latent Dirichlet Allocation* to perform [Topic Modeling](#) which is very helpful in classifying text specially when domain expertise is very basic.

size: support (0.1 – 0.2)
color: lift (1.4 – 1.6)



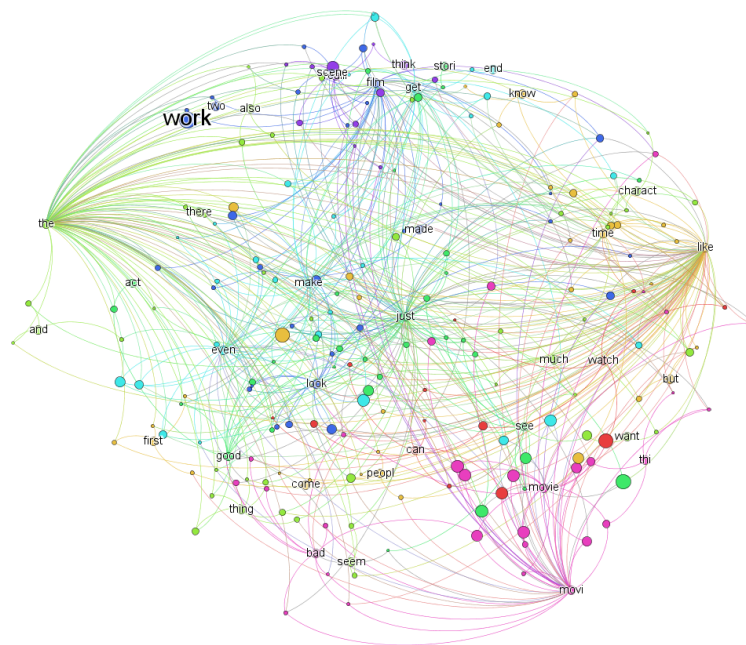


Figure 2: Text Network Analysis using Gephi

6 Decision Tree Learning

6.1 Most frequent words

```
findFreqTerms(keyword.term, lowfreq = 5000) # Which words are most frequent?
```

```
## [1] "act"      "actor"    "also"     "and"      "bad"      "best"     "better"
## [8] "but"      "can"      "charact"  "come"     "end"      "even"     "ever"
## [15] "film"     "find"     "first"    "get"      "give"     "good"     "great"
## [22] "just"     "know"     "life"     "like"     "littl"    "look"     "love"
## [29] "made"     "make"     "man"      "mani"     "movie"    "movi"     "much"
## [36] "never"    "one"      "peopl"    "plai"     "plot"     "realli"   "sai"
## [43] "scene"    "see"      "seem"     "seen"     "show"     "still"    "stori"
## [50] "take"     "the"      "there"    "thi"      "thing"    "think"    "time"
## [57] "two"      "wai"      "want"     "watch"    "well"     "will"     "work"
## [64] "year"
```

6.2 Bit-vector weighting

```
keyword.binary <- weightBin(keyword.term) # Applies binary transformation
keyword.binary <- as.data.frame.matrix(keyword.binary) # Converts to data frame
keyword.binary <- cbind(keyword.binary, sentiment = train$sentiment) # Combines response
keyword.binary <- as.data.frame(sapply(keyword.binary,
                                       FUN = as.factor)) # Converts columns to categorical
rownames(keyword.binary) <- row.names(train) # Applies unique ID
```

6.3 Divide train and testing set

```
set.seed(14344) # For reproducibility
split <- sample(seq_len(nrow(keyword.binary)), # Counts records from 1 to n
               size = 0.8 * nrow(keyword.binary), # Computes for 80% for training set
               replace = FALSE)

popcorn.train <- keyword.binary[split, ] # Creates training set
popcorn.test <- keyword.binary[-split, ] # Created testing set
```

6.4 Build tree

```
require(rpart)
require(rattle)
popcorn.tree <- rpart(formula = sentiment ~ ., # tests all predictors against response
                     data = popcorn.train, # From the training set
                     method = "class", # Tells model it is a classification tree
                     parms = list(split = "information"), # Uses information gain
                     model = T) # Retains model information
```

6.5 Visualize tree

```
# Prints nodes and leaves
```

```
print(popcorn.tree)
```

```
## n= 20000
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##  1) root 20000 9992 Pos (0.4996000 0.5004000)
##    2) bad=1 4583 1149 Neg (0.7492909 0.2507091) *
##    3) bad=0 15417 6558 Pos (0.4253746 0.5746254)
##      6) great=0 11498 5535 Pos (0.4813881 0.5186119)
##      12) love=0 9065 4362 Neg (0.5188086 0.4811914)
##        24) noth=1 983  276 Neg (0.7192269 0.2807731) *
##        25) noth=0 8082 3996 Pos (0.4944321 0.5055679)
##          50) best=0 6758 3213 Neg (0.5245635 0.4754365)
##            100) plot=1 1204  419 Neg (0.6519934 0.3480066) *
##            101) plot=0 5554 2760 Pos (0.4969391 0.5030609)
##              202) just=1 1814  765 Neg (0.5782800 0.4217200) *
##              203) just=0 3740 1711 Pos (0.4574866 0.5425134) *
##                51) best=1 1324  451 Pos (0.3406344 0.6593656) *
##          13) love=1 2433  832 Pos (0.3419647 0.6580353) *
##        7) great=1 3919 1023 Pos (0.2610360 0.7389640) *
```

```
layout(matrix(c(1,2,3,4), nrow = 1, ncol = 2, byrow = TRUE),
        widths=c(2.5,2))
```

```
# Plots the model
```

```
fancyRpartPlot(model = popcorn.tree,
               main = NULL,
               sub = NULL)
```

```
# Plots variable importance
```

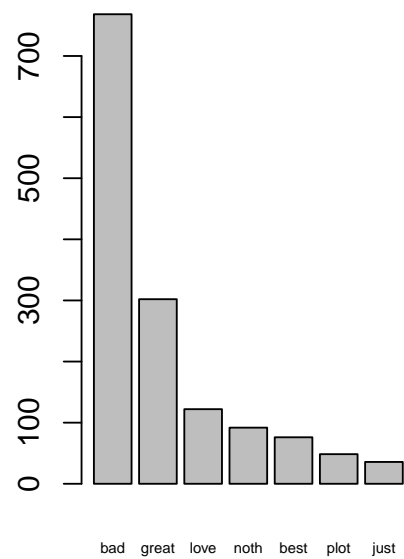
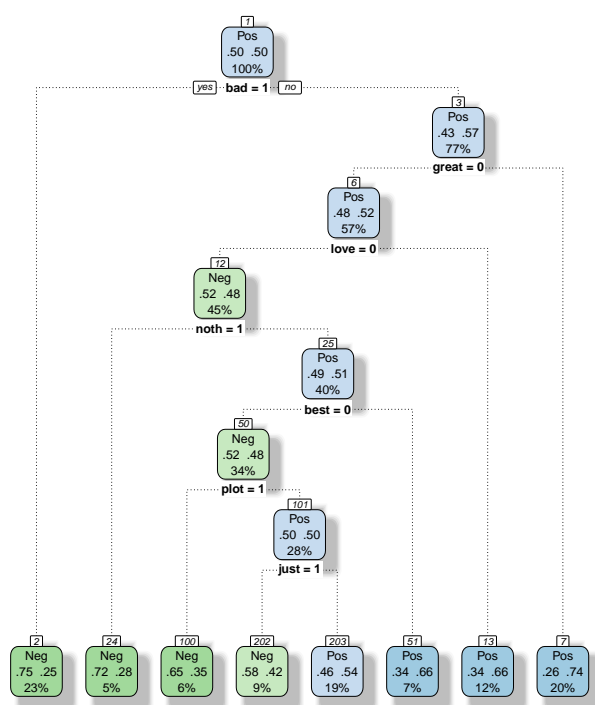
```
barplot(popcorn.tree$variable.importance,
        cex.names = 0.5,
        sub = "Variable Importance")
```

```
popcorn.tree$variable.importance
```

```
##      bad      great      love      noth      best      plot      just
## 768.29715 302.00945 122.08563  91.86011  76.09715  48.46538  35.75477
```

```
par(mfrow=c(1,1)) # Resets graphic device layout
```

6.6 Complexity parameter



Variable Importance

Figure 3:

```
printcp(popcorn.tree) # Prints CP
```

```
##
## Classification tree:
## rpart(formula = sentiment ~ ., data = popcorn.train, method = "class",
##       model = T, parms = list(split = "information"))
##
## Variables actually used in tree construction:
## [1] bad    best  great just  love  noth  plot
##
## Root node error: 9992/20000 = 0.4996
##
## n= 20000
##
##      CP nsplit rel error  xerror      xstd
## 1 0.228683      0  1.00000 1.01902 0.0070756
## 2 0.017064      1  0.77132 0.77132 0.0068882
## 3 0.015913      5  0.69496 0.73569 0.0068239
## 4 0.010000      7  0.66313 0.67394 0.0066887
```

Misclassification error = Root node error * standard error of lowest CP

$$0.4996 * 0.66783 = 0.3336479$$

More statistics are available in `summary(popcorn.tree)`.

6.7 Model evaluation

```
popcorn.prediction <- predict(object = popcorn.tree, # Tests model
                             newdata = popcorn.test, # with Testing set
                             type = "class") # Tells it is a classification prediction

# Build binary classification confusion matrix
popcorn.confusion <- table(Actual = popcorn.test$sentiment,
                          Predicted = popcorn.prediction)

tp <- popcorn.confusion[1,1] # True Positive
tn <- popcorn.confusion[2,2] # True Negative
fp <- popcorn.confusion[2,1] # False Positive
fn <- popcorn.confusion[1,2] # False Negative
n <- sum(popcorn.confusion) # Total records of testing set

popcorn.accuracy <- (tp + tn) / n # Accuracy rate = 66.78%
popcorn.error <- (fp + fn) / n # Error rate = 33.22%
popcorn.precision <- tp / (tp + fp) # Precision/Sensitivity = 69.50%
popcorn.recall <- tp / (tp + fn) # Recall/Specificity = 60.21%
popcorn.f1 <- 2 * popcorn.precision * popcorn.recall / (popcorn.precision + popcorn.recall)
popcorn.oddsratio <- (tp * tn) / (fp * fn) # Odds ratio = 4 folds
```

7 Citations

R Core Team (2015). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. .

Xie Y (2015). *knitr: A General-Purpose Package for Dynamic Report Generation in R*. R package version 1.10, .

Xie Y (2013). *Dynamic Documents with R and knitr*. Chapman and Hall/CRC, Boca Raton, Florida. ISBN 978-1482203530, .

Xie Y (2014). “knitr: A Comprehensive Tool for Reproducible Research in R.” In Stodden V, Leisch F and Peng RD (eds.), *Implementing Reproducible Computational Research*. Chapman and Hall/CRC. ISBN 978-1466561595, .

Hahsler M, Buchta C, Gruen B and Hornik K (2014). *arules: Mining Association Rules and Frequent Itemsets*. R package version 1.1-6, .

Hahsler M, Gruen B and Hornik K (2005). “arules - A Computational Environment for Mining Association Rules and Frequent Item Sets.” *Journal of Statistical Software*, 14(15), pp. 1-25. ISSN 1548-7660, .

Hahsler M and Chelluboina S (2014). *arulesViz: Visualizing Association Rules and Frequent Itemsets*. R package version 1.0-0, .

Boettiger C (2015). *knitcitations: Citations for knitr markdown files*. R package version 1.0.5, .

Therneau T, Atkinson B and Ripley B (2015). *rpart: Recursive Partitioning and Regression Trees*. R package version 4.1-9, .

Williams GJ (2011). *Data Mining with Rattle and R: The art of excavating data for knowledge discovery*, series Use R! Springer. .

Feinerer I and Hornik K (2014). *tm: Text Mining Package*. R package version 0.6, .

Feinerer I, Hornik K and Meyer D (2008). “Text Mining Infrastructure in R.” *Journal of Statistical Software*, 25(5), pp. 1-54. .

Bastian M., Heymann S., Jacomy M. (2009). Gephi: an open source software for exploring and manipulating networks. International AAAI Conference on Weblogs and Social Media.