

Contents

Motivation	1
Data set	1
Variables	2
Distribution	3
Regression	5
CART	5
Appendix	9
Data Mining	9
Data Scientist	11

Motivation

As I am going back to the Philippines to pursue further studies in Statistics, it intrigues me if Data Mining and Data Science are catching up. I am seeing some positions in jobsearch websites such as Jobstreet so as a data miner, I extracted the relevant job openings that are related to the key phrases: * **Data Mining**; and * **Data Scientist**.

These may look too specific but this is just a quick draft, anyway. Also, I did not include **Data Analyst** as this scopes a broader job scope diversity than the two mentioned not perform any intensive text extraction from Jobstreet.

Warning: The result of the models should not be used to provide recommendations as data the is collected using a convenience sample without performing accuracy tests, only k-fold cross validations against the training set when CART is used.

Data set

The data is collected manually by searching for relevant job openings active today, 22 May, 2015. I have an assumption that that the data set is relatively small, and so less than 30 positions returned. Pre-processing is done externally, in Excel, to remove currency prefix, *i.e.* PHP and text in experience, etc.

```
library(RCurl)
jobstreet <- getURL("https://raw.githubusercontent.com/foxyreign/Adhocs/master/Jobstreet.csv", ssl.verifi
writeLines(jobstreet, "Jobstreet.csv")
df <- read.csv('Jobstreet.csv', head=T, sep=",") # Load dataset
df <- na.omit(df) # Exclude missing data

summary(df) # Summarize
```

```
## Expected.Salary      Experience      Education
## Min.   : 11000      Min.    : 0.000      Min.   :1.000
## 1st Qu.: 19000      1st Qu.: 2.000      1st Qu.:2.000
## Median : 28000      Median : 4.000      Median :2.000
## Mean   : 35016      Mean    : 4.919      Mean    :2.179
## 3rd Qu.: 40000      3rd Qu.: 8.000      3rd Qu.:2.000
## Max.   :130000      Max.    :20.000      Max.    :4.000
##
##              Specialization      Position
## IT-Software      :30      Data Mining   :72
## -                 :17      Data Scientist:51
## IT-Network/Sys/DB Admin:14
## Actuarial/Statistics :12
## Banking/Financial   : 8
## Electronics         : 8
## (Other)             :34
```

As mentioned, there are only approximately 120 job applicants which applied for these two grouped positions. Since the data does not mention if an applicant applied for more than one position, I assume that these are distinct records of applicants per position and/or position group, Data Mining and Data Scientist.

Variables

1. **Expected.Salary** - numerical. The expected salary of each applicant based on their profile.
2. **Experience** - ordinal but treated as numerical for easier interpretation in the later algorithms used. This is the years of work experience of the applicant.
3. **Education** - categorical; not used in the models because of extreme unbalance in proportions. This is labelled as:
 - 1 - Secondary School
 - 2 - Bachelor Degree
 - 3 - Post Graduate Diploma
 - 4 - Professional Degree
3. **Specialization** - categorical; not used in this analysis.
4. **Position** - categorical. Data Mining or Data Scientist
5. **Education.Group** - categorical. Additional variable to bin the years of experience.

```
# Categorize education variable
df$Education <- factor(df$Education, levels = c(1,2,3,4),
                      labels=c("Secondary Sch", "Bach Degree",
                                "Post Grad Dip", "Prof Degree"))

# Bin years of experience
df$Experience.Group <- ifelse(df$Experience < 3, "3 Years",
                             ifelse(df$Experience < 5, "5 Years",
                                     ifelse(df$Experience < 10, "10 Years", "+10 Years")))
df$Experience.Group <- factor(df$Experience.Group,
                             levels=c("3 Years", "5 Years", "10 Years", "+10 Years"))

# Drop variables
df <- df[, !(colnames(df) %in% c("Education", "Specialization"))]
```

```
# Subsets positions
mining <- subset(df, Position == "Data Mining")
scientist <- subset(df, Position == "Data Scientist")
```

Distribution

As expected, Data Scientists have a higher expected salary although this is so dispersed that even if I compare these two using a t-test assuming heterodastic distribution, there is a significant difference between the averages expected salaries of the two positions.

```
require(ggplot2)
require(scales)

# Boxplot
ggplot(df, aes(x=factor(0), y=Expected.Salary, fill=Experience.Group)) +
  facet_wrap(~Position) + geom_boxplot() + xlab(NULL) +
  scale_y_continuous(labels = comma) +
  theme(axis.title.x=element_blank(),
        axis.text.x=element_blank(),
        axis.ticks.x=element_blank(),
        legend.position="bottom")
```

```
# T-test
t.test(Expected.Salary ~ Position, paired = FALSE, data = df)
```

```
##
##  Welch Two Sample t-test
##
## data:  Expected.Salary by Position
## t = -3.3801, df = 68.611, p-value = 0.001199
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -24086.501 -6205.983
## sample estimates:
##      mean in group Data Mining mean in group Data Scientist
##                28736.11                43882.35
```

```
# Median expected salaries of Data Mining vs Data Scientist
c(median(mining$Expected.Salary), median(scientist$Expected.Salary))
```

```
## [1] 25000 30000
```

Come on fellow data enthusiasts, you should do better than this! The difference of their medians is just 5,000 PHP. In my honest, these center values are way below based on the prospective demand of shortage of these people in the next 10 years.

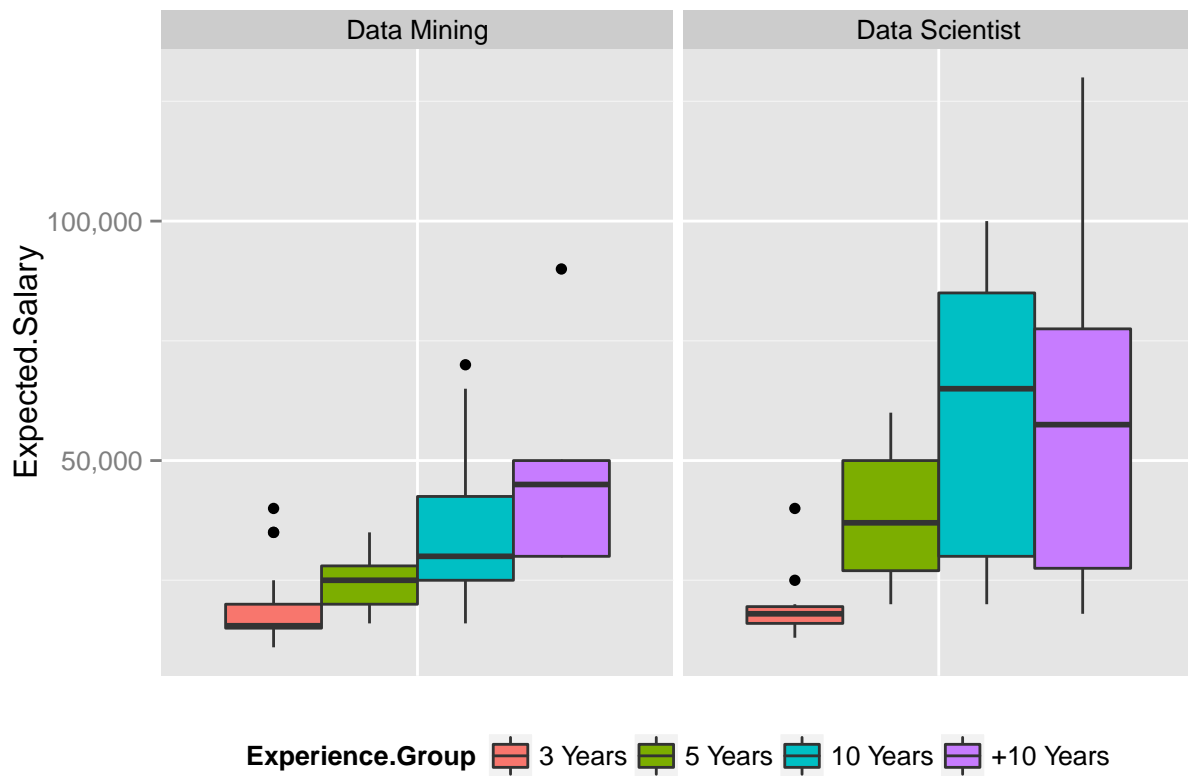


Figure 1: Distribution of Expected Salaries

Regression

The intercept is not included in the model because I want to see the contrast between Data Mining and Data Scientist although I already computed it beforehand. Besides, though the linear regression model shows significant value $r_{adj}^2 > 0.80, p < 0.05$ but when doing diagnostics, linear approach is not appropriate because the data is not random and depicts a funnel shape based on their errors.

The regression output coefficients are interpreted as follows:

$$y = \beta_0(12,934.9) + \beta_1(3,336.3) + \beta_2$$

```
# Estimate coefficients of linear regression model
summary(lm(Expected.Salary ~ Experience + Position-1, data=df))

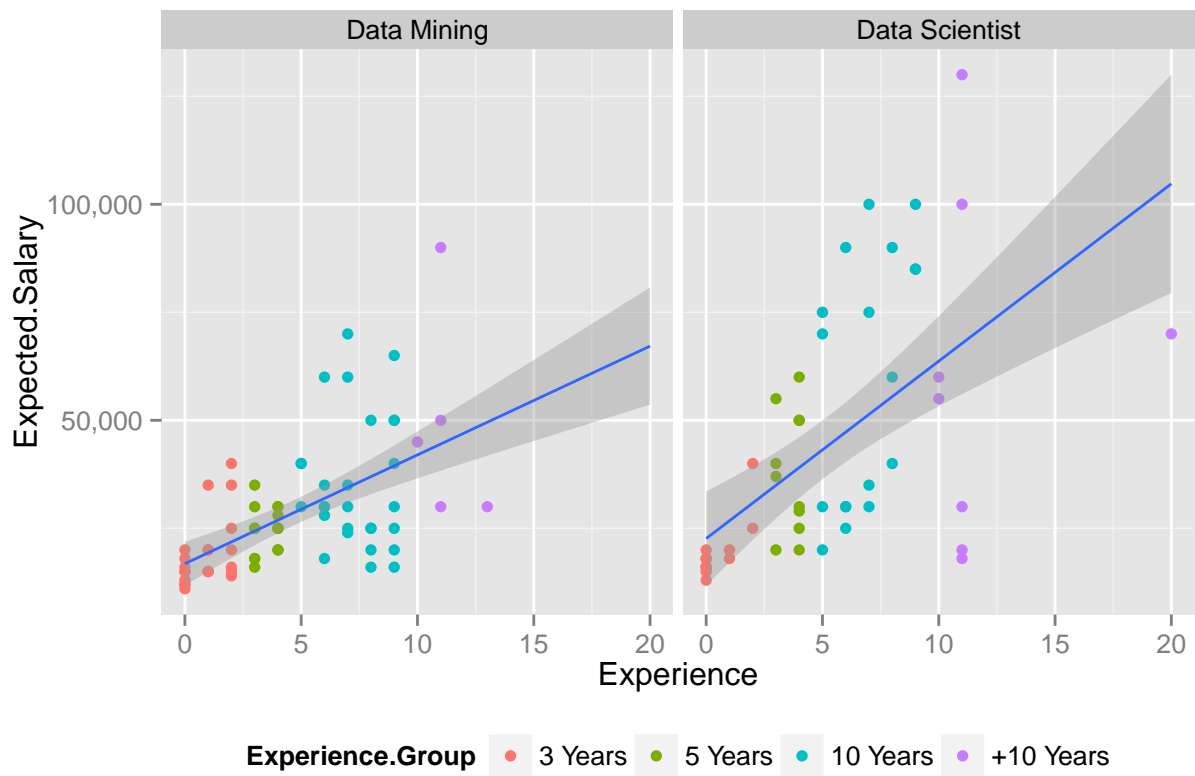
##
## Call:
## lm(formula = Expected.Salary ~ Experience + Position - 1, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -45312 -11123  -1280    6877   66688
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## Experience          3336.3      446.3   7.476 1.38e-11 ***
## PositionData Mining  12934.9     3024.3   4.277 3.83e-05 ***
## PositionData Scientist 26612.0     3455.8   7.701 4.27e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 18350 on 120 degrees of freedom
## Multiple R-squared:  0.8136, Adjusted R-squared:  0.809
## F-statistic: 174.6 on 3 and 120 DF,  p-value: < 2.2e-16

# Scatter plot
ggplot(df, aes(x=Experience, y=Expected.Salary)) +
  geom_point(aes(col=Experience.Group)) +
  facet_wrap(~Position) +
  scale_y_continuous(labels = comma) +
  stat_smooth(method="lm", fullrange = T) +
  theme(legend.position="bottom")

# Diagnose LM
par(mfrow=c(1,2))
plot(lm(Expected.Salary ~ Experience + Position-1, data=df), c(1,2))
```

CART

Information Gain is used to divide the nodes based on weighted average entropy as linear regression does not do well with the data set. Of course, years of experience is more influential than the position.



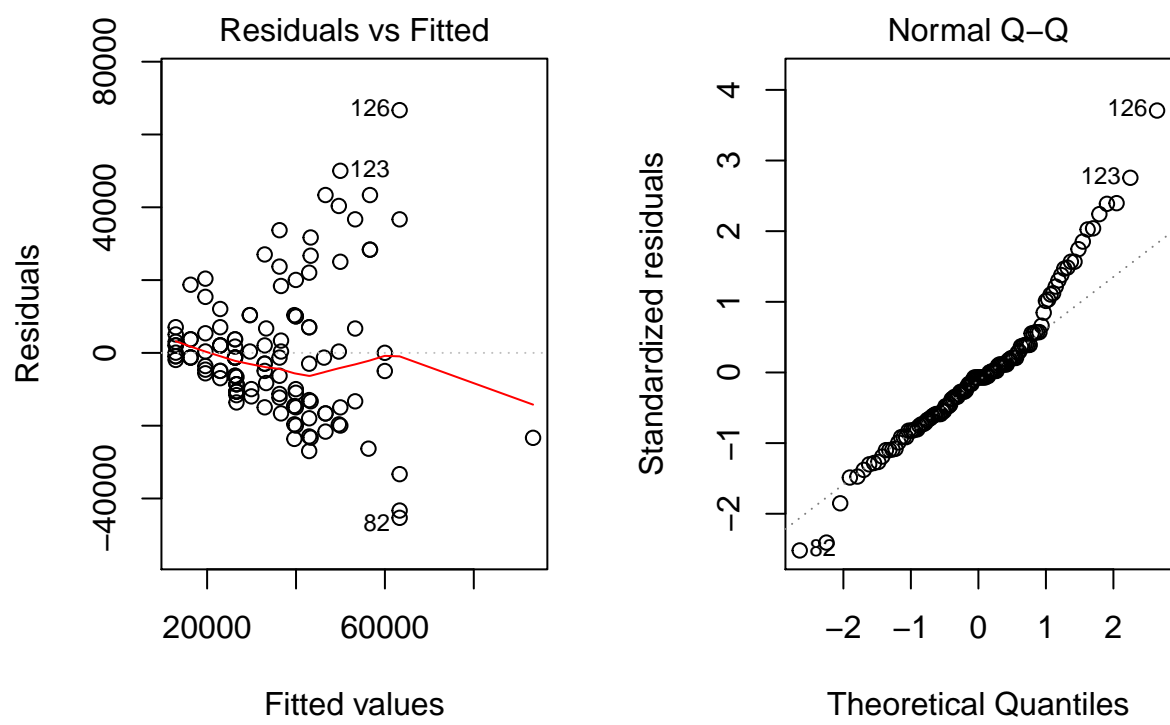


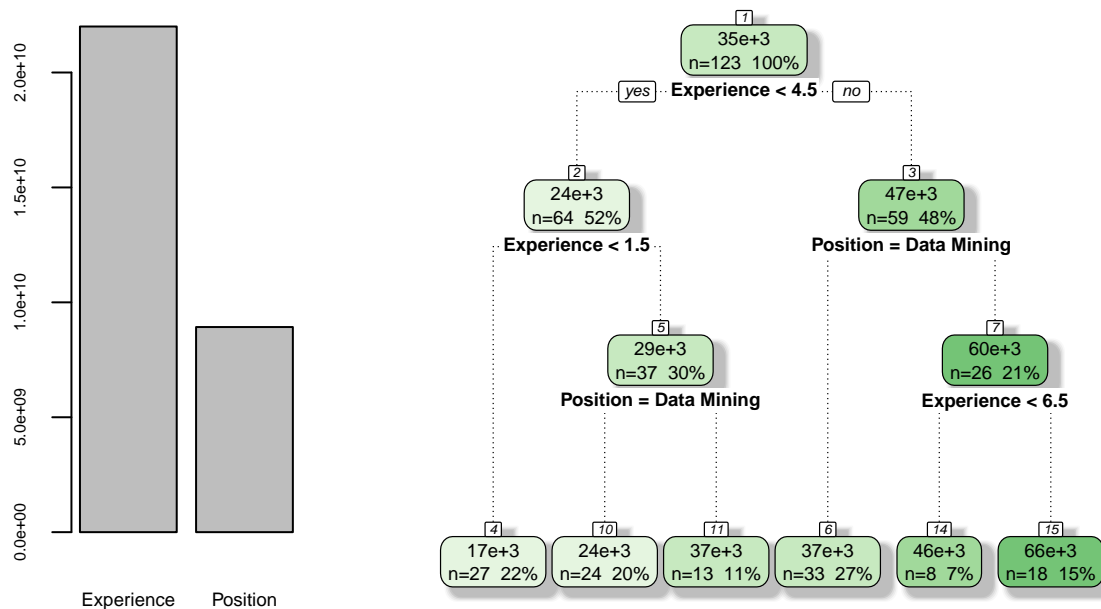
Figure 3: Regression and Diagnostics

Looking at the estimated salaries from the printed tree, applicants who have years of experience lower than 1.5 are approximately expecting 17,000 PHP. While does that applied for Data Mining jobs with 6.5 years of experience are expecting 66,000 pesos.

```
require(rpart)
require(rattle)

cart <- rpart(formula = Expected.Salary ~ Experience + Position,
              data = df,
              parms = list(split = "information"), # Uses information gain
              model = T) # Retains model information

# Plot tree
layout(matrix(c(1,2,3,4), nrow = 1, ncol = 2, byrow = TRUE), widths=c(1.5,2.5))
barplot(cart$variable.importance,
        cex.names = 0.6, cex.axis = 0.5,
        sub = "Variable Importance")
fancyRpartPlot(cart, main=NULL, sub=NULL)
```



```
# Estimates
print(cart); printcp(cart)
```

```
## n= 123
```



```

##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 123 66101970000 35016.26
##    2) Experience< 4.5 64 7326938000 23781.25
##      4) Experience< 1.5 27 520666700 16888.89 *
##      5) Experience>=1.5 37 4587676000 28810.81
##        10) Position=Data Mining 24 1055625000 24375.00 *
##        11) Position=Data Scientist 13 2188000000 37000.00 *
##    3) Experience>=4.5 59 41933560000 47203.39
##      6) Position=Data Mining 33 9501333000 37333.33 *
##      7) Position=Data Scientist 26 25137120000 59730.77
##        14) Experience< 6.5 8 5237500000 46250.00 *
##        15) Experience>=6.5 18 17799610000 65722.22 *

##
## Regression tree:
## rpart(formula = Expected.Salary ~ Experience + Position, data = df,
##       model = T, parms = list(split = "information"))
##
## Variables actually used in tree construction:
## [1] Experience Position
##
## Root node error: 6.6102e+10/123 = 537414370
##
## n= 123
##
##      CP nsplit rel error  xerror   xstd
## 1 0.254780      0  1.00000 1.00888 0.19245
## 2 0.110361      1  0.74522 0.81057 0.14709
## 3 0.033563      2  0.63486 0.77285 0.13292
## 4 0.031769      3  0.60130 0.76228 0.13052
## 5 0.020333      4  0.56953 0.74207 0.12969
## 6 0.010000      5  0.54919 0.70258 0.12366

```

Again, fellow data miners and data scientists, ask for more! You do not realize your worth with the current demand of people who can understand data.

Appendix

Data Mining

	Expected.Salary	Experience	Position	Experience.Group
## 1	60000	6	Data Mining	10 Years
## 2	16000	9	Data Mining	10 Years
## 3	18000	3	Data Mining	5 Years
## 4	20000	0	Data Mining	3 Years
## 5	20000	4	Data Mining	5 Years
## 6	20000	8	Data Mining	10 Years
## 7	25000	4	Data Mining	5 Years
## 8	30000	3	Data Mining	5 Years

## 9	30000	5 Data Mining	10 Years
## 10	35000	6 Data Mining	10 Years
## 11	40000	2 Data Mining	3 Years
## 12	40000	5 Data Mining	10 Years
## 13	40000	5 Data Mining	10 Years
## 14	40000	9 Data Mining	10 Years
## 15	50000	8 Data Mining	10 Years
## 16	50000	9 Data Mining	10 Years
## 17	50000	11 Data Mining	+10 Years
## 18	60000	7 Data Mining	10 Years
## 19	65000	9 Data Mining	10 Years
## 20	70000	7 Data Mining	10 Years
## 21	12000	0 Data Mining	3 Years
## 22	13000	0 Data Mining	3 Years
## 23	15000	0 Data Mining	3 Years
## 24	15000	1 Data Mining	3 Years
## 25	16000	0 Data Mining	3 Years
## 26	20000	1 Data Mining	3 Years
## 27	20000	1 Data Mining	3 Years
## 28	20000	2 Data Mining	3 Years
## 29	20000	4 Data Mining	5 Years
## 30	25000	2 Data Mining	3 Years
## 31	25000	4 Data Mining	5 Years
## 32	25000	4 Data Mining	5 Years
## 33	25000	4 Data Mining	5 Years
## 34	25000	8 Data Mining	10 Years
## 35	25000	8 Data Mining	10 Years
## 36	28000	4 Data Mining	5 Years
## 37	28000	6 Data Mining	10 Years
## 38	30000	4 Data Mining	5 Years
## 39	30000	6 Data Mining	10 Years
## 40	30000	9 Data Mining	10 Years
## 41	30000	13 Data Mining	+10 Years
## 42	35000	1 Data Mining	3 Years
## 43	35000	2 Data Mining	3 Years
## 44	35000	3 Data Mining	5 Years
## 45	35000	7 Data Mining	10 Years
## 46	45000	10 Data Mining	+10 Years
## 47	50000	9 Data Mining	10 Years
## 50	11000	0 Data Mining	3 Years
## 51	12000	0 Data Mining	3 Years
## 52	14000	2 Data Mining	3 Years
## 53	15000	0 Data Mining	3 Years
## 54	15000	1 Data Mining	3 Years
## 55	15000	1 Data Mining	3 Years
## 56	15000	2 Data Mining	3 Years
## 57	16000	2 Data Mining	3 Years
## 58	16000	3 Data Mining	5 Years
## 59	16000	8 Data Mining	10 Years
## 60	18000	0 Data Mining	3 Years
## 61	18000	3 Data Mining	5 Years
## 62	18000	6 Data Mining	10 Years
## 63	20000	9 Data Mining	10 Years
## 64	24000	7 Data Mining	10 Years

## 65	25000	3 Data Mining	5 Years
## 66	25000	3 Data Mining	5 Years
## 67	25000	3 Data Mining	5 Years
## 68	25000	7 Data Mining	10 Years
## 69	25000	9 Data Mining	10 Years
## 70	30000	4 Data Mining	5 Years
## 71	30000	6 Data Mining	10 Years
## 72	30000	7 Data Mining	10 Years
## 73	30000	11 Data Mining	+10 Years
## 74	90000	11 Data Mining	+10 Years

Data Scientist

##	Expected.Salary	Experience	Position	Experience.Group
## 75	13000	0 Data Scientist	3 Years	
## 76	15000	0 Data Scientist	3 Years	
## 77	16000	0 Data Scientist	3 Years	
## 78	16000	0 Data Scientist	3 Years	
## 79	18000	0 Data Scientist	3 Years	
## 80	18000	0 Data Scientist	3 Years	
## 81	18000	1 Data Scientist	3 Years	
## 82	18000	11 Data Scientist	+10 Years	
## 83	20000	0 Data Scientist	3 Years	
## 84	20000	1 Data Scientist	3 Years	
## 85	20000	11 Data Scientist	+10 Years	
## 86	29000	4 Data Scientist	5 Years	
## 87	30000	5 Data Scientist	10 Years	
## 88	30000	7 Data Scientist	10 Years	
## 90	16000	0 Data Scientist	3 Years	
## 91	16000	0 Data Scientist	3 Years	
## 92	18000	0 Data Scientist	3 Years	
## 93	20000	3 Data Scientist	5 Years	
## 94	20000	4 Data Scientist	5 Years	
## 95	20000	5 Data Scientist	10 Years	
## 96	25000	2 Data Scientist	3 Years	
## 97	25000	4 Data Scientist	5 Years	
## 98	25000	6 Data Scientist	10 Years	
## 99	30000	4 Data Scientist	5 Years	
## 100	30000	6 Data Scientist	10 Years	
## 101	30000	6 Data Scientist	10 Years	
## 102	30000	11 Data Scientist	+10 Years	
## 103	35000	7 Data Scientist	10 Years	
## 104	37000	3 Data Scientist	5 Years	
## 105	40000	2 Data Scientist	3 Years	
## 106	40000	3 Data Scientist	5 Years	
## 107	40000	8 Data Scientist	10 Years	
## 108	50000	4 Data Scientist	5 Years	
## 109	50000	4 Data Scientist	5 Years	
## 110	55000	3 Data Scientist	5 Years	
## 111	55000	10 Data Scientist	+10 Years	
## 112	60000	4 Data Scientist	5 Years	
## 113	60000	8 Data Scientist	10 Years	
## 114	60000	10 Data Scientist	+10 Years	

## 115	70000	5 Data Scientist	10 Years
## 116	70000	20 Data Scientist	+10 Years
## 117	75000	5 Data Scientist	10 Years
## 118	75000	7 Data Scientist	10 Years
## 119	85000	9 Data Scientist	10 Years
## 120	85000	9 Data Scientist	10 Years
## 121	90000	6 Data Scientist	10 Years
## 122	90000	8 Data Scientist	10 Years
## 123	100000	7 Data Scientist	10 Years
## 124	100000	9 Data Scientist	10 Years
## 125	100000	11 Data Scientist	+10 Years
## 126	130000	11 Data Scientist	+10 Years